

IRIS FLOWER CLASSIFICATION

The Iris flower dataset comprises three distinct species: Setosa, Versicolor, and Virginica. These species can be differentiated by their sepal and petal measurements. The objective is to utilize the Iris dataset for the development of a model capable of classifying Iris flowers into their respective species by analyzing their sepal and petal measurements.

Steps follwed

- Import The Libraries
- Data cleaning and preparation
- Visualization
- Split the data into Training and Testing
- Model building
- Conclusion

Import The Libraries

In [1]:

```
1 import os
2 import numpy as np
3 import pandas as pd
4 import seaborn as sns
5 import matplotlib.pyplot as plt
6 %matplotlib inline
7 sns.set(style="white", color_codes =True)
8
9 import warnings
10 warnings.filterwarnings('ignore')
```

In [2]:

```
1 df = pd.read_csv('IRIS.csv')
2 df.head(100)
```

Out[2]:

| | sepal_length | sepal_width | petal_length | petal_width | species |
|-----|--------------|-------------|--------------|-------------|-----------------|
| 0 | 5.1 | 3.5 | 1.4 | 0.2 | Iris-setosa |
| 1 | 4.9 | 3.0 | 1.4 | 0.2 | Iris-setosa |
| 2 | 4.7 | 3.2 | 1.3 | 0.2 | Iris-setosa |
| 3 | 4.6 | 3.1 | 1.5 | 0.2 | Iris-setosa |
| 4 | 5.0 | 3.6 | 1.4 | 0.2 | Iris-setosa |
| ... | ... | ... | ... | ... | ... |
| 95 | 5.7 | 3.0 | 4.2 | 1.2 | Iris-versicolor |
| 96 | 5.7 | 2.9 | 4.2 | 1.3 | Iris-versicolor |
| 97 | 6.2 | 2.9 | 4.3 | 1.3 | Iris-versicolor |
| 98 | 5.1 | 2.5 | 3.0 | 1.1 | Iris-versicolor |
| 99 | 5.7 | 2.8 | 4.1 | 1.3 | Iris-versicolor |

100 rows × 5 columns

In [3]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  -
0   sepal_length    150 non-null    float64
1   sepal_width     150 non-null    float64
2   petal_length    150 non-null    float64
3   petal_width     150 non-null    float64
4   species         150 non-null    object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

In [4]:

```
1 df.head()
```

Out[4]:

| | sepal_length | sepal_width | petal_length | petal_width | species |
|---|--------------|-------------|--------------|-------------|-------------|
| 0 | 5.1 | 3.5 | 1.4 | 0.2 | Iris-setosa |
| 1 | 4.9 | 3.0 | 1.4 | 0.2 | Iris-setosa |
| 2 | 4.7 | 3.2 | 1.3 | 0.2 | Iris-setosa |
| 3 | 4.6 | 3.1 | 1.5 | 0.2 | Iris-setosa |
| 4 | 5.0 | 3.6 | 1.4 | 0.2 | Iris-setosa |

Data cleaning and preparation

- Checking and removing duplicate values: In this step, we identify and remove any duplicate entries in the dataset to ensure that each observation is unique. Duplicate values can distort the analysis and lead to inaccurate results.
- Check for missing values and treat them: Missing values are instances where data is not recorded for certain variables. It is essential to identify and handle these missing values appropriately. Common methods for handling missing data include imputation (replacing missing values with estimated values) or deletion (removing rows or columns with missing data).
- Check for outliers and treat them: Outliers are extreme values that deviate significantly from the rest of the data. Outliers can skew the analysis and affect the model's performance. Identifying and treating outliers can involve methods like capping (replacing extreme values with a predetermined limit) or transformation (applying mathematical functions to normalize the data).

Check for missing values

In [5]:

```
1 df.isnull().sum()
```

Out[5]:

```
sepal_length    0
sepal_width     0
petal_length    0
petal_width     0
species         0
dtype: int64
```

In [6]:

```
1 df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   sepal_length    150 non-null   float64
1   sepal_width     150 non-null   float64
2   petal_length    150 non-null   float64
3   petal_width     150 non-null   float64
4   species         150 non-null   object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

In [7]:

```
1 df.shape
```

Out[7]:

(150, 5)

In [8]:

```
1 df.sample(5)
```


Out[8]:

| | sepal_length | sepal_width | petal_length | petal_width | species |
|-----|--------------|-------------|--------------|-------------|----------------|
| 8 | 4.4 | 2.9 | 1.4 | 0.2 | Iris-setosa |
| 120 | 6.9 | 3.2 | 5.7 | 2.3 | Iris-virginica |
| 15 | 5.7 | 4.4 | 1.5 | 0.4 | Iris-setosa |
| 40 | 5.0 | 3.5 | 1.3 | 0.3 | Iris-setosa |
| 35 | 5.0 | 3.2 | 1.2 | 0.2 | Iris-setosa |

Visualization


In [9]:

```
1 col1 = 'sepal_length'  
2 df[col1].hist()  
3 plt.suptitle(col1)  
4 plt.show()
```



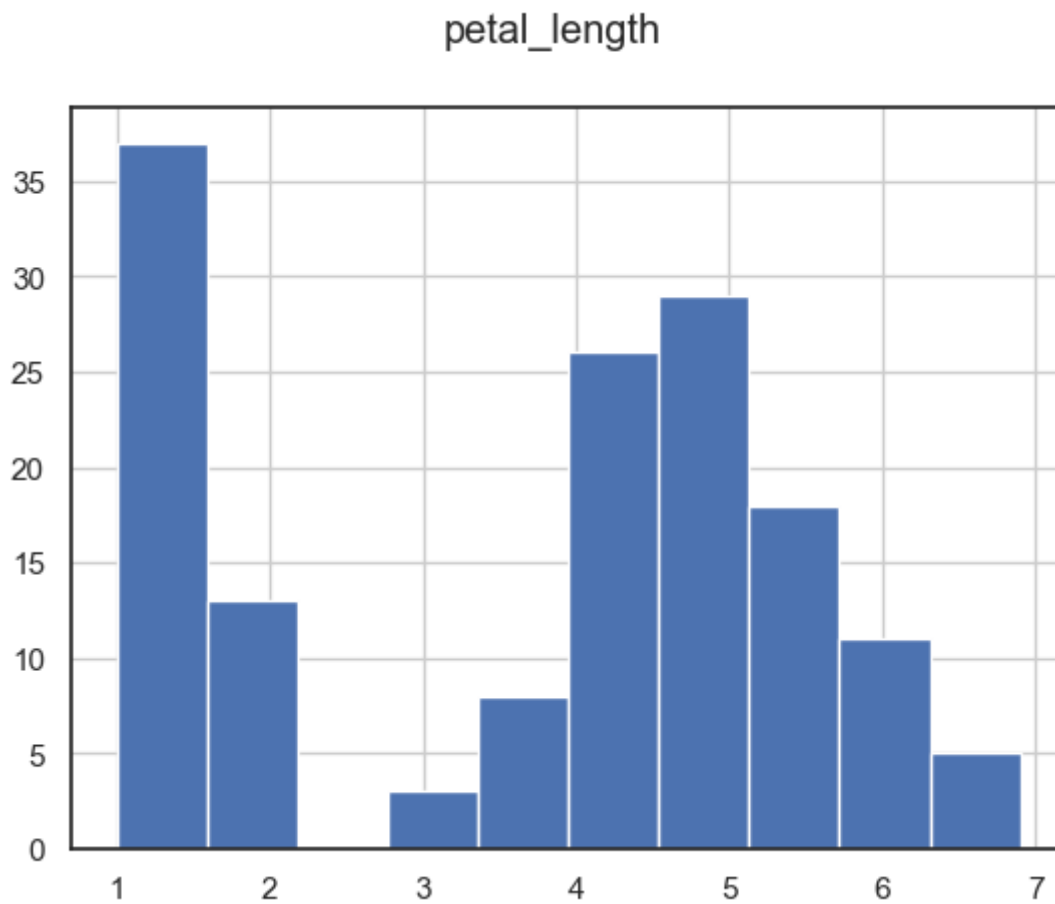
In [10]:

```
1 col2 = 'sepal_width'
2 df[col2].hist()
3 plt.suptitle(col2)
4 plt.show()
```




In [11]:

```
1 col3 = 'petal_length'  
2 df[col3].hist()  
3 plt.suptitle(col3)  
4 plt.show()
```



In [12]:


```
1 col4 = 'petal_width'
2 df[col4].hist()
3 plt.suptitle(col4)
4 plt.show()
```



Relationship between columns and species


In [13]:

```
1 col1 = 'sepal_length'
2 sns.relplot(x=col1, y='species', hue='species', data=df)
3 plt.title(col1, y=0.5)
4 plt.show()
```



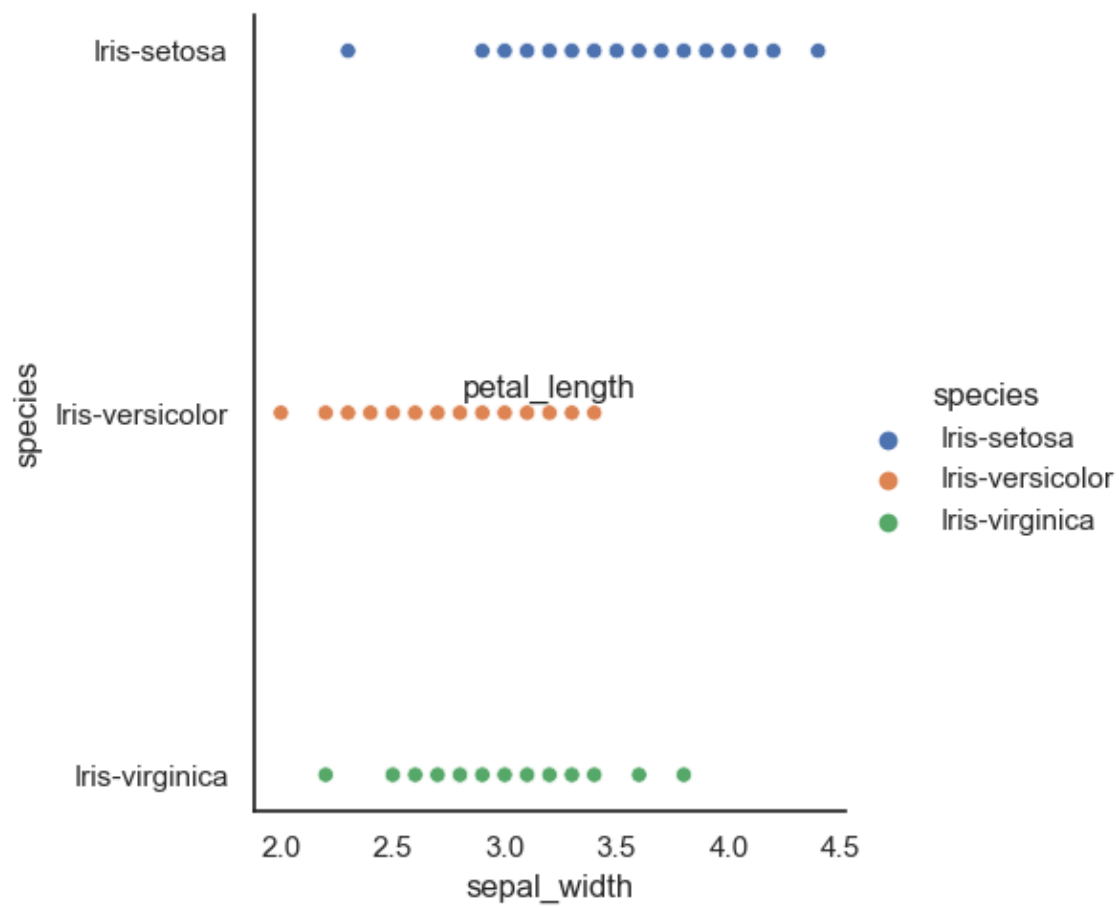
In [14]:

```
1 col2 = 'sepal_width'
2 sns.relplot(x=col2, y='species', hue='species', data=df)
3 plt.title(col2, y=0.5)
4 plt.show()
```




In [15]:

```
1 col3 = 'petal_length'  
2 sns.relplot(x=col2, y='species', hue='species', data=df)  
3 plt.title(col3, y=0.5)  
4 plt.show()
```



In [16]:

```
1 col4 = 'petal_width'
2 sns.relplot(x=col4, y='species', hue='species', data=df)
3 plt.title(col2, y=0.5)
4 plt.show()
```

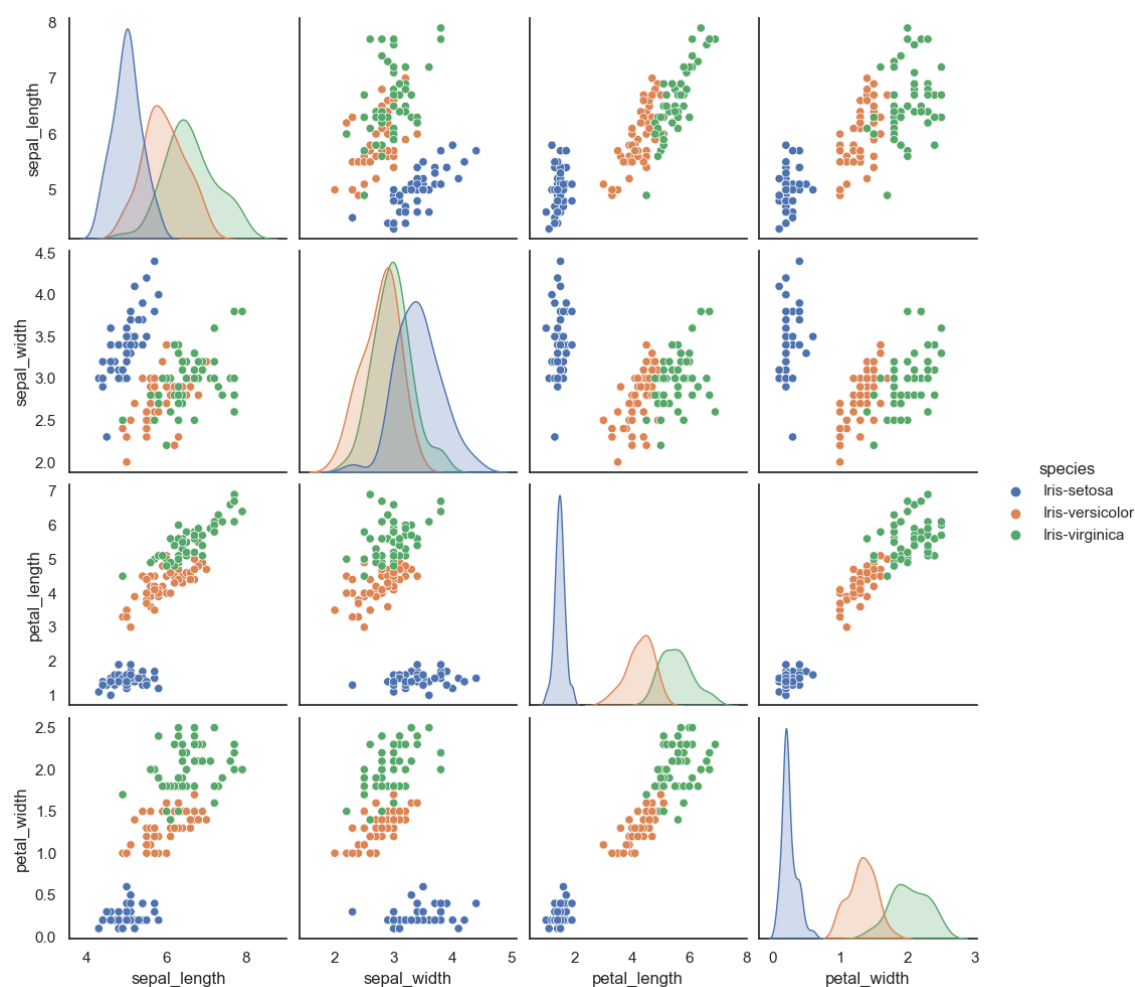


In [17]:

```
1 sns.pairplot(df,hue='species')
```

Out[17]:

<seaborn.axisgrid.PairGrid at 0x1736e1683d0>



Split the data into Training and Testing

In [18]:

```
1 from sklearn.preprocessing import LabelEncoder  
2 le =LabelEncoder()
```

In [19]:

```
1 le = le.fit_transform(df['species'])
2 df.head()
```

Out[19]:

| | sepal_length | sepal_width | petal_length | petal_width | species |
|---|--------------|-------------|--------------|-------------|-------------|
| 0 | 5.1 | 3.5 | 1.4 | 0.2 | Iris-setosa |
| 1 | 4.9 | 3.0 | 1.4 | 0.2 | Iris-setosa |
| 2 | 4.7 | 3.2 | 1.3 | 0.2 | Iris-setosa |
| 3 | 4.6 | 3.1 | 1.5 | 0.2 | Iris-setosa |
| 4 | 5.0 | 3.6 | 1.4 | 0.2 | Iris-setosa |

In [20]:

```
1 from sklearn.model_selection import train_test_split
2 df_train, df_test = train_test_split(df, test_size = 0.25)
```

In [21]:

```
1 df_train.shape
```

Out[21]:

(112, 5)

In [22]:

```
1 df_test.shape
```

Out[22]:

(38, 5)

Model building

In [23]:

```
1 x_train = df_train.drop(columns=["species"]).values
```

In [24]:

```
1 x_train
```

Out[24]:

```
array([[7.2, 3.6, 6.1, 2.5],
       [5.7, 2.9, 4.2, 1.3],
       [6. , 2.9, 4.5, 1.5],
       [5. , 3.5, 1.3, 0.3],
       [6.5, 3. , 5.8, 2.2],
       [5.4, 3.7, 1.5, 0.2],
       [5.6, 3. , 4.1, 1.3],
       [5. , 3. , 1.6, 0.2],
       [6.2, 2.9, 4.3, 1.3],
       [6.3, 2.9, 5.6, 1.8],
       [5.7, 2.8, 4.5, 1.3],
       [4.8, 3. , 1.4, 0.3],
       [7.1, 3. , 5.9, 2.1],
       [7.3, 2.9, 6.3, 1.8],
       [4.9, 3.1, 1.5, 0.1],
       [5.8, 2.7, 3.9, 1.2],
       [5.4, 3.9, 1.7, 0.4],
       [6.7, 3.1, 4.7, 1.5]])
```

In [25]:

```
1 y_train =df_train["species"].values
2
```

In [26]:

```
1 train, test= train_test_split(df,test_size= 0.25)
2 print(train.shape)
3 print(test.shape)
```

```
(112, 5)
```

```
(38, 5)
```

In [27]:

```
1 x_test = df_test.drop(columns=['species']).values
2 y_test = df_test['species'].values
```

In [28]:

```
1 x_test.shape
```

Out[28]:

```
(38, 4)
```

Logistic Regression

In [29]:

```

1
2 from sklearn.linear_model import LogisticRegression
3 from sklearn.metrics import accuracy_score

```

In [30]:

```

1 model = LogisticRegression(max_iter=1000)

```

In [31]:

```

1 model.fit(x_train,y_train)

```

Out[31]:

```
LogisticRegression(max_iter=1000)
```

In [32]:

```

1 prediction = model.predict(x_test)
2 print('Accuracy:',accuracy_score(prediction,y_test))

```

```
Accuracy: 0.9736842105263158
```

Confusion matrix

In [33]:

```

1 from sklearn.metrics import confusion_matrix,classification_report
2 confusion_mat = confusion_matrix(y_test,prediction)
3 print("Confusion matrix: \n",confusion_mat)
4 print(classification_report(y_test,prediction))

```

Confusion matrix:

```

[[12  0  0]
 [ 0 11  0]
 [ 0  1 14]]

```

| | precision | recall | f1-score | support |
|-----------------|-----------|--------|----------|---------|
| Iris-setosa | 1.00 | 1.00 | 1.00 | 12 |
| Iris-versicolor | 0.92 | 1.00 | 0.96 | 11 |
| Iris-virginica | 1.00 | 0.93 | 0.97 | 15 |
| accuracy | | | 0.97 | 38 |
| macro avg | 0.97 | 0.98 | 0.97 | 38 |
| weighted avg | 0.98 | 0.97 | 0.97 | 38 |

KNN Neighbors

In [34]:

```
1 from sklearn.neighbors import KNeighborsClassifier
2 model2 = KNeighborsClassifier(n_neighbors=5)
3 model2.fit(x_train,y_train)
4 y_pred2 = model2.predict(x_test)
5 from sklearn.metrics import accuracy_score
6 print("'Accuracy Score:", accuracy_score(y_test,y_pred2))
```

'Accuracy Score: 0.9736842105263158

Decision tree

In [35]:

```
1 from sklearn import tree
```

In [36]:

```
1 from sklearn.tree import DecisionTreeClassifier
2
3 dt_model = DecisionTreeClassifier()
4 dt_model.fit(x_train, y_train)
```

Out[36]:

DecisionTreeClassifier()

In [37]:

```
1 from sklearn.metrics import accuracy_score
2
3 prediction_dt = dt_model.predict(x_test)
4 accuracy_dt = accuracy_score (y_test,prediction_dt)
```

In [38]:

```
1 accuracy_dt
```

Out[38]:

0.9736842105263158

In [39]:

```

1
2
3 results = pd.DataFrame({
4     'model': ['Logistic Regression', 'KNN', 'Decision Tree'],
5     'Score': [0.97, 0.97, 0.97]
6 })
7
8 result_df = results.sort_values(by='Score', ascending=False)
9 result_df = result_df.set_index('Score')
10 result_df.head(5)

```

Out[39]:

| | model |
|-------|---------------------|
| Score | |
| 0.97 | Logistic Regression |
| 0.97 | KNN |
| 0.97 | Decision Tree |

In [40]:

```

1 from sklearn.model_selection import cross_val_score
2 scores= cross_val_score(dt_model,x_train, y_train, scoring='neg_mean_squared_error',
3 rmse_scores = np.sqrt(-scores)
4 rmse_scores

```

Out[40]:

```
array([nan, nan, nan, nan, nan, nan, nan, nan, nan, nan])
```

In [41]:

```
1 y_test
```

Out[41]:

```

array(['Iris-virginica', 'Iris-setosa', 'Iris-versicolor',
      'Iris-virginica', 'Iris-virginica', 'Iris-setosa', 'Iris-setosa',
      'Iris-setosa', 'Iris-virginica', 'Iris-versicolor',
      'Iris-virginica', 'Iris-versicolor', 'Iris-virginica',
      'Iris-virginica', 'Iris-virginica', 'Iris-versicolor',
      'Iris-virginica', 'Iris-virginica', 'Iris-versicolor',
      'Iris-versicolor', 'Iris-versicolor', 'Iris-setosa',
      'Iris-versicolor', 'Iris-setosa', 'Iris-setosa', 'Iris-virginica',
      'Iris-virginica', 'Iris-virginica', 'Iris-setosa', 'Iris-setosa',
      'Iris-versicolor', 'Iris-setosa', 'Iris-setosa', 'Iris-versicolor',
      'Iris-versicolor', 'Iris-virginica', 'Iris-setosa',
      'Iris-versicolor'], dtype=object)

```

In [42]:

```
1 prediction_dt
```

Out[42]:

```
array(['Iris-virginica', 'Iris-setosa', 'Iris-versicolor',  
      'Iris-virginica', 'Iris-virginica', 'Iris-setosa', 'Iris-setosa',  
      'Iris-setosa', 'Iris-virginica', 'Iris-versicolor',  
      'Iris-virginica', 'Iris-versicolor', 'Iris-virginica',  
      'Iris-virginica', 'Iris-versicolor', 'Iris-versicolor',  
      'Iris-virginica', 'Iris-virginica', 'Iris-virginica',  
      'Iris-versicolor', 'Iris-versicolor', 'Iris-setosa',  
      'Iris-versicolor', 'Iris-setosa', 'Iris-setosa', 'Iris-virginica',  
      'Iris-virginica', 'Iris-virginica', 'Iris-setosa', 'Iris-setosa',  
      'Iris-versicolor', 'Iris-setosa', 'Iris-setosa', 'Iris-versicolor',  
      'Iris-versicolor', 'Iris-virginica', 'Iris-setosa',  
      'Iris-versicolor'], dtype=object)
```

creating the category

In [43]:

```
1 category = ['setosa', 'versicolor', 'virginica']
```

In [44]:

```
1 data= 5.7,3.4,4.2,1.1
```

In [45]:

```
1 data_array = np.array([data])  
2 data_array
```

Out[45]:

```
array([[5.7, 3.4, 4.2, 1.1]])
```

In [46]:

```
1 predict =dt_model.predict(data_array)
```

In [47]:

```
1 print(predict[0])
```

Iris-versicolor

Input from user

In [48]:

```
1 sepal_length = float(input("Enter Sepal Length :"))
2 sepal_width = float(input("Enter Sepal Width:"))
3 petal_length = float(input("Enter Petal Length :"))
4 petal_width = float(input("Enter Petal Width :"))
5
6
7 # convert user input into Numpy array
8 input_data = np.array([[sepal_length, sepal_width, petal_length, petal_width]])
9
10 # use the trained model to predict the species of flower
11 predicted_species = dt_model.predict(input_data)
12
13 # Display the predicted species to the user
14 print( "Predicted Species:", predicted_species[0])
```

```
Enter Sepal Length :5
Enter Sepal Width:3.6
Enter Petal Length :1
Enter Petal Width :1.2
Predicted Species: Iris-versicolor
```

Conclusion

throughout our analysis, we observed that the various species of Iris flowers display discernible characteristics concerning sepal and petal measurements. By employing machine learning algorithms such as Logistic Regression, Decision Trees, and K-Nearest Neighbors (KNN) classification, we successfully constructed models capable of precisely classifying Iris flowers into their respective species based on these measurements. The dataset's features, including Sepal Length (in centimeters), Sepal Width (in centimeters), Petal Length (in centimeters), and Petal Width (in centimeters), were crucial in enabling the accurate differentiation of Iris flower species. Our results highlight the effectiveness of machine learning techniques in automating species classification based on morphological attributes, contributing to the fields of botany, ecology, and beyond.

In []:

1